

Navigation

June 1, 2020

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.
```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```
In [2]: from unityagents import UnityEnvironment
import numpy as np

# please do not modify the line below
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
Number of Brains: 1
```

```
Number of External Brains : 1
```

```
Lesson number : 0
```

```
Reset Parameters :
```

```
Unity brain name: BananaBrain
```

```
Number of Visual Observations (per agent): 0
```

```
Vector Observation space type: continuous
```

```
Vector Observation space size (per agent): 37
```

```

Number of stacked Vector Observation: 1
Vector Action space type: discrete
Vector Action space size (per agent): 4
Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)
        state_size = len(state)
        print('States have length:', state_size)

```

```

Number of agents: 1
Number of actions: 4
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755      1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
States have length: 37

```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that in this coding environment, you will not be able to watch the agent while it is training, and you should set `train_mode=True` to restart the environment.

```
In [5]: env_info = env.reset(train_mode=True)[brain_name] # reset the environment
        state = env_info.vector_observations[0]           # get the current state
        score = 0                                         # initialize the score
        while True:
            action = np.random.randint(action_size)      # select an action
            env_info = env.step(action)[brain_name]      # send the action to the environment
            next_state = env_info.vector_observations[0]  # get the next state
            reward = env_info.rewards[0]                 # get the reward
            done = env_info.local_done[0]                 # see if episode has finished
            score += reward                               # update the score
            state = next_state                           # roll over the state to next time step
            if done:                                     # exit loop if episode finished
                break

        print("Score: {}".format(score))
```

Score: 1.0

When finished, you can close the environment.

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

1.0.5 Step 5: Import the dependencies

```
In [13]: from collections import deque
         import matplotlib.pyplot as plt
         import random
         import torch

         %matplotlib inline
```

1.0.6 Step 6: Start Implementing

```
In [14]: from agent import Agent
```

1.0.7 Step 7: Deep Q Learning

```
In [15]: # Deep Q-Learning Function
```

```
def dqn(n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.995, train_mode=True, ckpt_path='saved_weights/weights.pth'):  
    """Deep Q-Learning.  
  
    Params  
    =====  
    n_episodes (int): maximum number of training episodes  
    max_t (int): maximum number of timesteps per episode  
    eps_start (float): starting value of epsilon, for epsilon-greedy action selection  
    eps_end (float): minimum value of epsilon  
    eps_decay (float): multiplicative factor (per episode) for decreasing epsilon  
    train_mode (bool): if 'True' set environment to training mode  
  
    """  
    scores = [] # list containing scores from each episode  
    scores_window = deque(maxlen=100) # last 100 scores  
    moving_avgs = [] # list of moving averages  
    eps = eps_start # initialize epsilon  
    for i_episode in range(1, n_episodes+1):  
        env_info = env.reset(train_mode=train_mode)[brain_name] # reset environment  
        state = env_info.vector_observations[0] # get current state  
        score = 0  
        for t in range(max_t):  
            action = agent.act(state, eps) # select an action  
            env_info = env.step(action)[brain_name] # send action to environment  
            next_state = env_info.vector_observations[0] # get next state  
            reward = env_info.rewards[0] # get reward  
            done = env_info.local_done[0] # see if episode has finished  
            agent.step(state, action, reward, next_state, done) # learning step  
            state = next_state  
            score += reward  
            if done:  
                break  
        scores_window.append(score) # save most recent score to window  
        scores.append(score) # save most recent score to total  
        moving_avg = np.mean(scores_window) # calculate moving average  
        moving_avgs.append(moving_avg) # save most recent moving average  
        eps = max(eps_end, eps_decay*eps) # decrease epsilon  
        print('\rEpisode {} \t Average Score: {:.2f}'.format(i_episode, moving_avg), end=' ')  
        if i_episode % 100 == 0:
```

```

        print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, moving_avg))
    if moving_avg >= 13.0:
        print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, moving_avg))
        if train_mode:
            torch.save(agent.qnetwork_local.state_dict(), ckpt_path)
        break
    return scores, moving_avgs

```

1.0.8 Best Performing Agent

- Standard DQN + replay buffer (no double, no dueling)

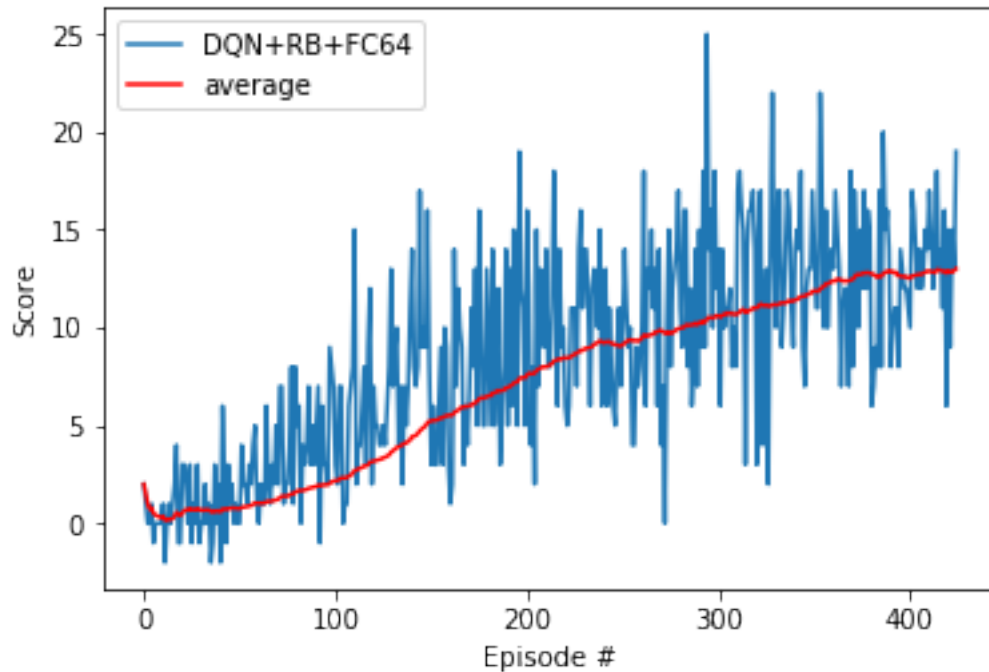
```

In [18]: # run the training loop
agent = Agent(state_size=state_size, action_size=action_size, seed=0, use_double=False,
              scores, avgs = dqn(n_episodes=600, eps_decay=0.98, eps_end=0.02, ckpt_path='saved_weights.pth'))

# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='DQN+RB+FC64')
plt.plot(np.arange(len(scores)), avgs, c='r', label='average')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left');
plt.show()

```

Episode 100	Average Score: 2.16	
Episode 200	Average Score: 7.44	
Episode 300	Average Score: 10.62	
Episode 400	Average Score: 12.56	
Episode 425	Average Score: 13.01	
Environment solved in 325 episodes!		Average Score: 13.01



1.0.9 Testing of Best Performing Agent

In [20]: *## Test the saved agent*

```
# initialize the agent
agent = Agent(state_size=state_size, action_size=action_size, seed=0)

# load the weights from file
checkpoint = 'saved_weights/final_weights.pth'
agent.qnetwork_local.load_state_dict(torch.load(checkpoint))

num_episodes = 10
scores = []
for i_episode in range(1,num_episodes+1):
    env_info = env.reset(train_mode=False)[brain_name] # reset the environment
    state = env_info.vector_observations[0] # get the current state
    score = 0 # initialize the score
    while True:
        action = agent.act(state, eps=0) # select an action
        env_info = env.step(action)[brain_name] # send the action to the environment
        next_state = env_info.vector_observations[0] # get the next state
        reward = env_info.rewards[0] # get the reward
        done = env_info.local_done[0] # see if episode has finished
        #agent.step(state, action, reward, next_state, done) # do the learning
```


```

        score += reward                                # update the score
        state = next_state                             # roll over the state to next time
        if done:                                       # exit loop if episode finished
            scores.append(score)
            print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores)))
            break

# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

Episode 1	Average Score: 19.00
Episode 2	Average Score: 16.50
Episode 3	Average Score: 17.33
Episode 4	Average Score: 16.00
Episode 5	Average Score: 13.00
Episode 6	Average Score: 13.33
Episode 7	Average Score: 13.29
Episode 8	Average Score: 13.50
Episode 9	Average Score: 13.56
Episode 10	Average Score: 13.60



1.0.10 Summary

```
In [23]: from IPython.display import Image
```

```
Image(filename='images/Results.png', width=800)
```

Out[23]:

Iteration	Agent	Model (FC1 units)	Eps Decay	Eps End	Result (# episodes)			
28	DQN+RB	64	0.98	0.02	200			
8	DQN+RB	64	0.95	0.03	216		Agent Legend DQN Deep Q-Network DDQN Double DQN Dueling Dueling DQN RB Replay Buffer	
29	DDQN+RB	64	0.98	0.02	231			
9	DDQN+RB+Dueling	64	0.98	0.02	232			
10	DQN+RB	64	0.98	0.02	245			
26	DQN+RB+Dueling	128	0.98	0.02	246			
19	DQN+RB+Dueling	128	0.985	0.015	248			
12	DQN+RB	64	0.99	0.01	263			
25	DDQN+RB	128	0.98	0.02	263			
20	DQN+RB	128	0.985	0.015	266			
31	DDQN+RB+Dueling	64	0.98	0.02	274			
22	DQN+RB	128	0.985	0.05	285			
21	DQN+RB	128	0.95	0.03	294			
18	DDQN+RB	128	0.985	0.015	297			
23	DQN+RB	128	0.985	0.01	312			
27	DDQN+RB+Dueling	128	0.98	0.02	312			
17	DDQN+RB+Dueling	128	0.985	0.015	317			
30	DQN+RB+Dueling	64	0.98	0.02	328			
11	DDQN+RB+Dueling	64	0.99	0.01	333			
24	DQN+RB	128	0.98	0.02	362			
2	DQN+RB	64	0.995	0.01	380			

```
In [ ]:
```